

# ECOM90025 Advanced Data Analysis - Pandas Basic

Zheng Fan

July 31, 2023

## 1 Basic commands for pandas

- Just like a library in R, “pandas” is a popular open-source data manipulation and analysis library. It provides easy-to-use data structures and functions for efficiently working with structured data.
- A 10-minuts to pandas, covering nearly everything we need.  
[https://pandas.pydata.org/docs/user\\_guide/10min.html](https://pandas.pydata.org/docs/user_guide/10min.html)
- Ask ChatGPT

To use pandas in your Python environment, you need to install it first using the following command:  
- remove ! in front to install

```
[1]: !pip install pandas
```

```
Requirement already satisfied: pandas in c:\anaconda3\lib\site-packages (1.2.4)
Requirement already satisfied: pytz>=2017.3 in c:\anaconda3\lib\site-packages
(from pandas) (2021.1)
Requirement already satisfied: numpy>=1.16.5 in c:\anaconda3\lib\site-packages
(from pandas) (1.20.1)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\anaconda3\lib\site-
packages (from pandas) (2.8.1)
Requirement already satisfied: six>=1.5 in c:\anaconda3\lib\site-packages (from
python-dateutil>=2.7.3->pandas) (1.15.0)
```

Once installed, you can import it in your Python scripts or Jupyter notebooks using:

```
[2]: import pandas as pd
```

- pd is therefore just a shortcut pandas, you can choose what you like.
- we can use it to make a simple data frame

```
[3]: df_sample = pd.DataFrame(
    {
        "Var_1": 1.0,
        "Var_2": pd.Categorical(["test", "train", "test", "train"]),
        "Var_3": "foo",
    }
)
print(df_sample)
```

	Var_1	Var_2	Var_3
0	1.0	test	foo
1	1.0	train	foo
2	1.0	test	foo
3	1.0	train	foo

We need to check the data type before doing any numerical exercise

- check python data type here: [https://www.w3schools.com/python/python\\_datatypes.asp](https://www.w3schools.com/python/python_datatypes.asp)

```
[4]: df_sample.dtypes
```

```
[4]: Var_1      float64
     Var_2      category
     Var_3      object
     dtype: object
```

Check data dimension

```
[5]: # the dimension of the data set
     df_sample.shape
```

```
[5]: (4, 3)
```

For illustration purpose, I will use a publicly available dataset from a online Monash python class

```
[6]: df = pd.read_csv("https://monashdatafluency.github.io/python-workshop-base/
     ↪modules/data/surveys.csv")
```

Viewing data

```
[7]: df.head(2)
```

```
[7]:   record_id  month  day  year  site_id species_id sex  hindfoot_length \
0         1      7   16  1977         2         NL   M             32.0
1         2      7   16  1977         3         NL   M             33.0

     weight
0      NaN
1      NaN
```

```
[8]: df.tail(2)
```

```
[8]:   record_id  month  day  year  site_id species_id sex  hindfoot_length \
35547    35548     12   31  2002         7         D0   M             36.0
35548    35549     12   31  2002         5        NaN  NaN             NaN
```

```

weight
35547    51.0
35548    NaN

```

```
[9]: df.columns
```

```
[9]: Index(['record_id', 'month', 'day', 'year', 'site_id', 'species_id', 'sex',
          'hindfoot_length', 'weight'],
          dtype='object')
```

Shows a quick statistic summary of your data:

```
[10]: df.describe()
```

```
[10]:
```

	record_id	month	day	year	site_id \
count	35549.000000	35549.000000	35549.000000	35549.000000	35549.000000
mean	17775.000000	6.474022	16.105966	1990.475231	11.397001
std	10262.256696	3.396583	8.256691	7.493355	6.799406
min	1.000000	1.000000	1.000000	1977.000000	1.000000
25%	8888.000000	4.000000	9.000000	1984.000000	5.000000
50%	17775.000000	6.000000	16.000000	1990.000000	11.000000
75%	26662.000000	9.000000	23.000000	1997.000000	17.000000
max	35549.000000	12.000000	31.000000	2002.000000	24.000000

  

	hindfoot_length	weight
count	31438.000000	32283.000000
mean	29.287932	42.672428
std	9.564759	36.631259
min	2.000000	4.000000
25%	21.000000	20.000000
50%	32.000000	37.000000
75%	36.000000	48.000000
max	70.000000	280.000000

Selecting a single column, which yields a Series

```
[11]: df["weight"]
```

```
[11]: 0      NaN
      1      NaN
      2      NaN
      3      NaN
      4      NaN
      ...
      35544  NaN
      35545  NaN
```

```

35546    14.0
35547    51.0
35548     NaN
Name: weight, Length: 35549, dtype: float64

```

### Get some rows

```
[12]: df[0:3]
```

```

[12]:   record_id  month  day  year  site_id  species_id  sex  hindfoot_length  \
0         1      7   16  1977         2          NL   M             32.0
1         2      7   16  1977         3          NL   M             33.0
2         3      7   16  1977         2          DM   F             37.0

      weight
0      NaN
1      NaN
2      NaN

```

### Selecting by label

```
[13]: df.loc[:, ["day", "site_id"]]
```

```

[13]:   day  site_id
0    16         2
1    16         3
2    16         2
3    16         7
4    16         3
...
35544  31        15
35545  31        15
35546  31        10
35547  31         7
35548  31         5

[35549 rows x 2 columns]

```

### Performing a descriptive statistic:

```
[14]: df.mean()
```

```

[14]: record_id    17775.000000
      month         6.474022
      day         16.105966
      year        1990.475231

```

```
site_id          11.397001
hindfoot_length  29.287932
weight           42.672428
dtype: float64
```

```
[15]: df.mean(1) # calculate row mean
```

```
[15]: 0          339.166667
      1          339.666667
      2          340.333333
      3          341.166667
      4          340.500000
      ...
      35544      7521.000000
      35545      7521.200000
      35546      5375.857143
      35547      5383.857143
      35548      7519.800000
      Length: 35549, dtype: float64
```

```
[16]: df.month.mean()
```

```
[16]: 6.474021772764353
```